

JY-CRPG-BENCH: A LONG-HORIZON BENCHMARK FOR FRONTIER MODELS IN AN OPEN-WORLD RPG

Anonymous authors

Paper under double-blind review

ABSTRACT

A long-horizon task requires a model to perceive its environment, decide for itself what is worth doing, and hold one plan across hours of dependent decisions. We present JY-CRPG-BENCH, which places a frontier model inside Heroes of Jin Yong, a 1996 open-world role-playing game run under DOS emulation. Its ending, recovering the fourteen books of the Jin Yong novels, takes a first-time human player tens of hours. The game requires reading 320×200 isometric pixel art and 16-pixel Traditional Chinese, knowledge of the novels its world is built from, and a keyboard as the only tool. It has no quest log, marker or score, so the model must maintain its own map and goals, and it must win turn-based fights against the game AI. Progress is scored by milestones of the game itself and, once a model finishes, by completion time. Under a 60-minute budget, one of 11 frontier models plays past the opening of the game, and none gains experience or holds a book.

1 INTRODUCTION

Frontier models solve competition mathematics and write working compilers, yet the tasks they complete reliably are still measured in hours of human work (Kwa et al., 2025). A long horizon is a property of the task, the number of dependent steps its completion requires (Chen et al., 2026), and the benchmarks that measure it have so far run in text, as business simulations (Backlund & Petersson, 2025; Fan et al., 2026) or interactive fiction (Phan et al., 2025). A game adds the requirement they leave out: the model has to recover its own state from raw perception, because nothing in the environment reports it. Existing game benchmarks each remove one part of that problem, by rendering the state as text, removing language, or drawing on a widely known setting (Paglieri et al., 2025; ARC Prize Foundation, 2026; Ju et al., 2026), so the full combination is untested.

We present JY-CRPG-BENCH, built on Heroes of Jin Yong (金庸群俠傳), a 1996 Chinese role-playing game that we run as its original DOS binary under emulation. Its ending requires recovering the fourteen books of the Jin Yong novels from an open world, a task that takes a first-time human player tens of hours. The game demands several capabilities of a frontier model at once, and Figure 1 shows its world. The model must read 320×200 pixel art drawn in a 45° isometric projection, with Traditional Chinese in 16-pixel glyphs and no key that moves straight up the screen. It must bring knowledge the game never states. The world is assembled from the fourteen novels, so who a character is, which faction to join and which storyline leads to a book are known to a reader of the novels and written nowhere on screen. The only actuator is a keyboard, exposed through an API that the model calls as a tool. The game keeps no quest log, draws no marker and reports no score, so the model must maintain its own model of the world and decide for itself what is worth doing. ARC-AGI-3 places the same goal-setting requirement at the centre of agentic intelligence (ARC Prize Foundation, 2026). The fights of the game are turn-based on a grid against its AI, with position, range and inner force to manage. Together these form one long-horizon task, a persistent world with sparse and delayed reward over thousands of dependent steps. This is the setting of reinforcement learning, and here a frontier model plays it in context, with no training on the game. Under a 60-minute budget, one of 11 frontier models plays past the opening, and even that model is still at its first fight, with all fourteen books still to find.

Table 1 places JY-CRPG-BENCH among interactive agent benchmarks. The next section reviews game environments as agent benchmarks and the measurement problem that long horizons create.



Figure 1: The environment as the model receives it, upscaled three times without interpolation.

Section 3 details the environment, the interface, the session protocol and the measurement, and Section 4 reports which milestones the models reach and how they play.

Table 1: Interactive agent benchmarks. Observation is what the model receives before any upscaling, a self-directed environment supplies no task list, score or objective marker during play, and the last column is the record that verifies progress.

Benchmark	Environment	Observation	Language	Self-directed	Progress from
JY-CRPG-BENCH	one 1996 CRPG	320×200 px	Trad. Chinese	yes	memory, save, replay
BALROG	6 research games	text, render	English	no	env. reward
TextQuests	interactive fiction	text	English	yes	game state
ARC-AGI-3	abstract puzzles	64×64 grid	none	yes	env. score
VideoGameBench	1990s action games	native px	English	no	checkpoint images
GameWorld	34 browser games	browser	English	no	game state
MineExplorer	Minecraft	3D render	English	no	milestone rules

2 RELATED WORK

2.1 GAME ENVIRONMENTS AS AGENT BENCHMARKS

Games have been the standard setting for sequential decision-making since the Arcade Learning Environment (Bellemare et al., 2013), and language and vision-language agents have produced a second generation of game benchmarks. BALROG aggregates six reinforcement-learning environments under one protocol and reports a persistent gap between what a model knows about a game and what it does inside one (Paglieri et al., 2025). Imgame-Bench isolates how much of a game result belongs to the harness (Hu et al., 2026), and Orak spans twelve popular titles with configurable agentic modules (Park et al., 2026). VideoGameBench is the nearest benchmark to ours, running real 1990s games from raw pixels under a no-scaffold rule (Zhang et al., 2025a), and Atari-GPT scores

multimodal models as low-level policies in Atari games (Waytowich et al., 2024). GameWorld standardises the action interface across 34 browser games and pairs every task with a state-verifiable metric (Ouyang et al., 2026). MineExplorer evaluates open-world exploration in Minecraft and finds that models handle single-hop tasks but degrade once hidden prerequisites must be coordinated over a longer trajectory (Ju et al., 2026).

The nearest precedents in horizon are the public Pokémon runs. Claude 3.7 Sonnet reached three gym badges from screen pixels behind a memory scaffold, where an earlier model had not left the first house (Anthropic, 2025). Gemini 2.5 Pro finished Pokémon Blue in 813 hours behind a harness that translated the RAM of the game into text, and an ablation without vision performed about as well (Comanici et al., 2025). The PokeAgent Challenge turned the setting into a speedrunning track scored by completion time over milestones (Karten et al., 2026a). PokeGym builds a long-horizon benchmark on the open-world game Pokémon Legends: Z-A with no access to game state (Zhang et al., 2026a), and StarBench one on the turn-based Honkai: Star Rail (Zhang et al., 2025b). Crafter established the single-environment, achievement-based design we follow (Hafner, 2022), and NetHack remains the deepest environment in use for learned agents (Küttler et al., 2020).

2.2 LONG-HORIZON EVALUATION AND ITS MEASUREMENT

Vending-Bench established long-term coherence as the property that separates long-horizon tasks from long chains of short ones (Backlund & Petersson, 2025), and E-Commerce Bench extends the setting to a year of deterministic business operation (Fan et al., 2026). Both keep perception out of the loop. TextQuests shares our self-directed, long-horizon framing while removing vision (Phan et al., 2025), and TextAtari stretches horizons to 10^5 steps over text renderings (Li et al., 2025). ARC-AGI-3 isolates adaptation to novel tasks by excluding language and external knowledge (ARC Prize Foundation, 2026), and AutumnBench probes world-model quality through environment-level queries (Warrier et al., 2025). UltraHorizon makes exploration under hidden rules the task itself (Luo et al., 2025), and Kapoor et al. (2026) observe that benchmark design has favoured tasks that are precisely specified, automatically graded and short, and call for open-world evaluation. Navigation from rendered images is a documented weakness: on mazes, one-shot navigation collapses for every model by ten cells a side (Jiang et al., 2026), and the failures are attributed to looping (Einarsson, 2025).

Long horizons also complicate measurement. The ABC checklist catalogues task-validity and outcome-validity failures across agentic benchmarks (Zhu et al., 2025). An audit of agent traces finds reward hacking in a majority of them on two suites (Shao et al., 2026), and a survey of the horizon gap finds that the field answers uninformative outcome signals by building denser step-level ones (Chen et al., 2026). State-based verification has precedent outside games, where τ -bench compares the final database state with an annotated goal (Yao et al., 2024).

3 JY-CRPG-BENCH

3.1 ENVIRONMENT

Heroes of Jin Yong, published by Heluo Studio (河洛工作室) in 1996, opens with the player in a small compound. The only standing instruction comes from the fiction: enter the world of the Jin Yong novels, join its factions, learn its martial arts, and recover the fourteen books that end the game. The world is two-tier, with many local scenes such as buildings, sects and caves attached to one world map of 480×480 tiles (scarsty & contributors, 2026). Each of 320 named characters carries a record of level, experience, health, stamina, skills, items, reputation and aptitude. Every scene carries an entry condition. Of the 84 scenes, 7 are open at the start: the home, five inns (客棧) and the house of the hermit (南賢). The conversation with the hermit opens all but 4 of the rest, and from there the model chooses the order of play.

The compound holds a searchable chest, one speaking character and one exit tile. Beyond it, the path south leads to the hermit. A cabinet beside him holds the compass (羅盤), the only source of coordinates the game offers, and his conversation opens the house of the first companion, who joins after a yes-or-no prompt at which the confirm key means no. Fights are turn-based on the scene grid. Each turn a character moves along the four axes and then strikes, uses an item, waits or rests. Skills draw on inner force (内力), and the first scripted fight begins with the party drained of it. A party

that goes elsewhere first meets storylines whose fights it cannot survive at level one. Some of these fights end the game when lost, leaving only the load menu, and others let the story continue.

Every objective, prompt and statistic reaches the model as Traditional Chinese rendered at sixteen pixels a line, a script that multimodal evaluation has largely neglected (Tam et al., 2025). Two further properties made the game suitable: it is menu-driven, so a model that deliberates between actions is not penalised as in a real-time game (cf. Zhang et al., 2025a), and its progression state is measurable per character. Appendix A states the licences of the benchmark code, the emulator and the game.

3.2 OBSERVATION AND ACTION SPACE

The model observes the raw framebuffer at 320×200 and acts by pressing keys. A scored session serves the native frame without upscaling, annotation, set-of-marks overlay or accessibility tree (cf. Xie et al., 2024). The frame combines several documented weaknesses of vision-language models. Low resolution reduces their accuracy (Niu et al., 2025), corrupted input reduces it further (Usama et al., 2025), and images far from the training distribution degrade recognition even of familiar categories (Lin et al., 2026). Text rendered as pixels is understood less well than the same text supplied as tokens (Liu et al., 2026), and precise spatial reading fails once primitives are small or adjacent (Rahmanzadehgervi et al., 2024; Wang et al., 2024). The action space is the DOS keyboard, with 130 key names over 100 distinct keys, and Figure 5 in Appendix B shows the few the game responds to. Movement runs along the four numpad diagonals, with the arrow keys aliased onto the same axes, so a model that presses up expecting to walk up the screen moves up and to the right. The view follows the character except near the edge of a scene, where the background holds still and the sprite walks across it, so a model that reads movement from the scrolling background loses count of its tiles there. Canopies and roofs are drawn over the character.

The single action presses a key, or a list of up to 100 keys in order, holding each for a requested number of emulated frames, ten by default. It returns once the screen has settled. The game has 30 frames to react, and the wait then ends after 9 frames without change, when the picture cycles as in an idle animation, or after 120 frames regardless. A black frame marks a scene change, and the action waits through it and settles once more, so the frame a model receives is the one to act on. There is no call that advances the game without input, since the game only moves on input, and the hold is counted in frames because the game samples its keyboard once per game-loop iteration. A scored session answers four of the calls in Table 2: the screen, the brief, the key names and the action. Three further calls save, list and restore emulator snapshots for development, and a scored session refuses them, so a run cannot be rewound from outside the game. The save and load menu of the game stays available to the model as to a human player, so a model that loses a fight can reload the last save the game wrote. Appendix C gives the reasons for the list action and for the other choices of the protocol.

Table 2: The control surface. A check marks a call the model can use in a scored session; the snapshot calls serve development.

Call	Meaning	Model can use
GET /api/screen	the current frame	✓
GET /api/help	the brief, in English or Chinese	✓
GET /api/keys	every accepted key name	✓
GET /api/slots	emulator snapshots on disk	×
POST /api/key {"key": "kp3", "hold": 10}	press one key, or several in order	✓
POST /api/save {"name": "..."}	take a snapshot	×
POST /api/load {"name": "..."}	restore a snapshot	×

3.3 HARNESS AND SESSION PROTOCOL

The benchmark supplies the brief and the control surface, and the model plays through an agent harness. Every model in this paper ran inside pi¹, a minimal coding agent installed as shipped and set to its high thinking level. Its four tools read, write and edit files and run a shell, with no web

¹<https://github.com/badlogic/pi-mono>

access, so no model could look up a walkthrough. The model drives the control surface with HTTP calls from that shell. The brief is the whole instruction the model receives. It covers the calls that constitute play, the rules of a run, the opening route as far as the compass, and a manual of the controls and systems of the game. It is published as `agents.md`, the project file `pi` reads at launch, and served in English and Chinese with the on-screen Traditional Chinese terms kept verbatim, so the language a model reasons in is not itself a handicap. An action returns only its metadata, so the model looks at the screen when it decides to, and the benchmark adds no memory scaffold, no map, no pathfinder and no access to the state of the game. A model keeps any memory between actions in files it writes in that shell, and a scored session receives no field that reports its progress until the run ends, so a model cannot query its own score. The leaderboard therefore ranks models under one fixed harness, since harness choice changes game results (Hu et al., 2026; Karten et al., 2026b) and on long-horizon tasks the variance a harness induces can exceed the variance between models (Zhang et al., 2026b).

A session is created under a declared name with a playtime, 60 minutes by default and up to twenty-four hours, and the clock starts at the first playable moment. The wall clock of every decision call is kept on the server, so played time and think time are measured for any client and shown beside a run without ranking it. Each session is a separate emulator process with a private copy of the game, supplied by the host, with its save slots reset to a new game, so no run inherits the progress of another. A session ends only when its budget lapses or the game is completed, and there is no inactivity cutoff. When the budget lapses the session renders its recording to a time-compressed video, computes its metrics, and appends itself to the public catalogue with the video and its keypress timeline, so the behaviour behind a score can be reviewed. Figure 2 shows the apparatus, and Appendix B lists the emulation and hosting parameters.




Figure 2: A frontier model in pi drives the control API of its own session. The recording feeds the public replay, and progress is read from the emulated machine, the save the game writes and the replay, none of which the control API exposes.

3.4 MEASURING PROGRESS

The game keeps its persistent state in one archive: a base block with the world position, the party roster and the shared bag, followed by 320 character records of 182 bytes and an item table. We validated that layout against the opening machine image and the shipped archives. The memory reader locates the array in the live machine by anchoring on one character name and confirming the two neighbouring records at the expected stride. The bag in front of that array is the working copy and changes the moment an item is taken. The item count, the compass and the books are read from it after every action and again from every save the game writes. A milestone, once observed,

stays credited, so an item picked up and later used still counts. Reading is passive: the benchmark serialises the machine and never loads state into it. The only keys it presses on its own belong to the save attempt described next. The attempt takes the game a few seconds of menu time that is not charged to the run, and it leaves the screen and the game state unchanged, since the menu is modal and the game advances only on input.

For the party roster and the world position the benchmark reads the save the game writes. The game offers its save menu on the world map alone, and the menu itself reveals whether the party is there, since it has six rows on the world map and four in a scene. An attempt opens the menu, reads its height, and saves only where six rows appear. Attempts recur every two minutes and once more before the budget lapses, each waiting for a two-second gap in the keys the model sends. The session is held while the attempt runs, so the model never sees the screen mid-save, and the time of the first successful save is the time of the crossing. Neither the archive nor anything read from it is reachable through the control surface, so a successful save is itself the evidence that the party has crossed onto the world map. The session record also keeps a fingerprint of the world map taken from the frame, which credits the crossing only for a session with no save on record.

Five events leave no trace in either record, because the game keeps them only on screen. The model speaks with the hermit, opens the item screen with the compass in it, enters a fight, loses it, or answers the prompt of a companion who asks to join. These are read from the published replay, since each event draws fixed art at a fixed screen position. The hermit has a portrait in the dialogue frame, and the compass adds a coordinate line to the item screen. A fight shows the card of the acting character, a defeat shows its banner (戰鬥失敗), and a recruitable character asks with a yes-or-no prompt. Every frame of the replay is matched against these panels at one frame per second by normalised cross-correlation, and a score above 0.9 counts as a hit. A prompt followed by a yes key in the keypress journal is a recruitment. The message the game centres on screen when an item enters the bag is matched along its row at the full frame rate of the video. It supplies the item milestone for the 6 sessions with no bag reading and agrees with the bag on all 36 sessions that have one. Appendix C reports the same check for the other readings.

Progress is summarised as 10 milestones, in the spirit of the achievements of Crafter (Hafner, 2022). The milestones fall into two phases that the game itself separates. The opening needs no fight. Its five milestones are the exit tile, the chest in the first room, the conversation with the hermit, the compass in his cabinet, and the companion who joins through dialogue, all reached by walking, reading and answering. The campaign begins with the first fight, which a party at level one can lose, and runs through the experience and levels that victories bring and the chains of items, factions and fights that lead to the books. Its five milestones are a fight entered, a fight fought to its end, experience gained, the second level, and the first of the fourteen books. The default budget covers the opening, and extended budgets cover the campaign. The milestones are listed in the order of the opening route, with the optional chest after the exit, and a run is credited with every milestone it reaches, so the count does not depend on the path taken.

The milestones extend to the ending of the game. The fourteen books² are required before the final sequence can be played, and the archive records every book held, so the count of books continues the scale beyond the last milestone. The first read that finds all fourteen records the completion time, which ranks the sessions that finish. For multi-hour budgets the release also records the scenes reached, battles survived, skills learned and books gathered, so a long run without a book still leaves an ordered record of its progress.

4 EVALUATION

We evaluate 11 models from 5 vendors at the 60-minute budget, every one run in pi at the high thinking level, with 40 sessions on record and 2 to 6 per model. A model is credited with every milestone any of its sessions reached, and a session created under a variant of a model name, a spelling or the thinking level, is listed under the model. The baseline is a seeded random policy that

²Item records 144 to 157 of the item table, named after the novels: 飛狐外傳、雪山飛狐、連城訣、天龍八部、射鵬英雄傳、白馬嘯西風、鹿鼎記、笑傲江湖、書劍恩仇錄、神鵰俠侶、俠客行、倚天屠龍記、碧血劍、鴛鴦刀. A book counts as held in the shared bag or carried by the player. With all fourteen and enough reputation, the endgame is a boss fight at a shrine on the far side of the world map (Bahamut forum contributors, 2023).

plays the same budget through the same control surface, drawing from 13 keys the brief teaches, weighted towards movement and confirmation, and never reading the screen (Mnih et al., 2015; Hafner, 2022; ARC Prize Foundation, 2026).

4.1 MAIN RESULTS

Of the 11 models, all 11 reach the world map, 9 pick up an item, 4 speak with the hermit, and one, gpt-6-astra, takes the compass, recruits a companion, enters a fight and fights it to the end. It reaches 7 of the 10 milestones across its 4 sessions. No model gains experience: every one of the 35 character records read remains at level 1 with 0 experience and 1 skill, as it was when the game began, so every milestone from experience onward remains unreachable. Figure 3 shows, for each model, the share of sessions that reached each milestone. Comparable benchmarks report a similar position: BALROG places its best model at 1.5% progression on NetHack (Paglieri et al., 2025), and VideoGameBench reports the best vision-language models completing 0.48% of its suite (Zhang et al., 2025a). Everything between the first fight and the fourteenth book remains as headroom.




Figure 3: Milestones per model, with rows ordered by the mean number of actions to the world map, fewest first, and the random baseline last. In the upper panel each marker shows the share of sessions in which the model reached the milestone, with the number of sessions at the right. The lower panel shows the actions to the world map across the sessions that crossed, on a log scale, with a light dot per session and a black dot for their mean. Each milestone is read from the record in brackets: the world map from a save written there or the world-map fingerprint (save, frame); the item from the bag count or the pickup message (memory, replay); the hermit from his portrait (replay); the compass from the bag or the item screen (memory, replay); the companion from the party size or a yes to the join prompt (save, replay); the fight from the card of the acting character (replay); the fight fought to the end from the defeat banner or experience gained (replay, memory); experience and level from the character record (memory); a book from the bag or the player (memory). A session whose replay shows no fight gained no experience and holds no book, since both require a won fight.

4.2 QUALITATIVE ANALYSIS

Three milestones separate the models: the chest, the hermit and the fight. Figure 4 shows one frame from each behaviour pattern behind these divisions. The random policy reaches 1 of the 10 milestones in 60 minutes, since in a small room a random walk searches the chest within the hour. Its recomputed key sequence matches the journal of the machine key for key, and no save succeeds during its hour, so the world map is the first milestone a random walk does not reach. The chest therefore separates models from one another but not from the baseline: 9 models searched it and 2 never did, and the milestone count summarises what a model achieved without assuming one path.



Figure 4: Frames from the replays, each with the strip naming the model, the action count and the clock: (a) the compass coordinates on the item screen; (b) the conversation with the hermit; (c) the join prompt of Duan Yu in an inn; (d) a fight; (e) the hermit naming the cabinet; (f) the yard behind the fence; (g) passing the gate of the hermit; (h) talking to an innkeeper; (i) the defeat banner.

Every model crosses onto the world map, in 30 of the 40 sessions, and 6 models cross in every session they played. The replays show where the actions go. The compound is left through a yard closed by a fence with one gate, and several sessions were stuck in that yard. glm-5.3-flash needed 1595 actions to cross in one session and reached the world map at minute 37 after long stretches of walking into the fence, against a median crossing of 86 actions. It never pressed enter or space in that session, so the chest in the first room, which opens on either, stayed shut. The count varies more within a model than between models, as the lower panel of Figure 3 shows. The crossings of glm-5.3-flash differ by a factor of 25 and the means of the models by a factor of 16, while gpt-6-astra crossed in 28 to 43 actions in all 4 sessions. The game itself times the crossing, since the first save written on the world map trails it by at most the two-minute cadence of the attempts. By that clock the 10 timed crossings fall between 9 and 41 minutes into the session, with a median of 26.

Beyond the fence, the hour is spent on the world map. claude-fable5.1 crossed at minute 13, then spent 1261 actions and 44 minutes on the world map without entering a scene, and passed

the gate of the hermit repeatedly. The entrance of a walled compound is one tile on one face, and every other face is solid. Two sessions that found an inn and a shop talked to the keeper repeatedly, although the replies do not advance the opening. `gpt-5.6-terra` sent lists of up to 100 keys in one action, 22 on average, so it looked at the screen once for every 22 keys it pressed, and its crossing is credited by the save alone.

The second division is the hermit, and it turns on reading his dialogue. This milestone and the fights and recruitment above it are read from the replay. `gemin-3.7-flash`, `gemin-3.8-flash`, `gpt-5.6-sol` and `gpt-6-astra` reach his conversation in 8 sessions, and the conversation is long. `gemin-3.8-flash` pressed through his 40 lines at 23 seconds each, so it took 15 minutes and ended at minute 54, when the hermit had named the cabinet and 6 minutes of budget remained. `gemin-3.7-flash` reached him at minute 37 and pressed 60 times in 1 minute. Neither searched the cabinet. `gpt-6-astra` reached him in every one of its 4 sessions, between minutes 11 and 19. The 3 other models that reached him did so in at most 2 of their 4 sessions. The compass-holding session that entered the most scenes, 11 in the hour, pressed through the conversation in 7 minutes at 10 seconds a line and searched the cabinet. It then opened the item screen 9 times to read its coordinates from the compass, the one instrument the brief recommends, and its 42 arrow-key presses all fell inside menus and prompts, never on the world map.

Beyond the compass only one model remains. All 4 sessions of `gpt-6-astra` read the compass between minutes 18 and 29, and from there they diverged. One met Duan Yu (段譽) at an inn at minute 46, answered yes at minute 52 when he asked to travel together, and was handed the manual of his art (六脈神劍譜). One entered a fight at minute 27 and lost it at minute 28, then opened the system menu, loaded the save the benchmark had written, and by minute 34 held the compass again. One entered a fight at minute 40 and sent its last key at minute 42 with the fight still open. The fight is the third division and the last, since no model reached the campaign milestones beyond it.

The budget is wall clock, so the compute a model spends on a decision is charged against the time it has left to play, an efficiency framing shared with ARC-AGI-3 and TextQuests (ARC Prize Foundation, 2026; Phan et al., 2025). The sessions took between 31 and 1994 actions, and Table 3 in Appendix B lists actions, keys per action and the time between actions for every model. A comparison of token cost across models must wait until more than one model plays past the opening. A model that has not reached the hermit has not begun the part of the game the milestones score, and at this budget only `gpt-6-astra` has, so a cost per milestone would rank a single model.

5 CONCLUSION

We presented JY-CRPG-BENCH, a long-horizon benchmark built on a 1996 open-world role-playing game that demands the full repertoire of a frontier model at once. The model must perceive isometric pixel art and Traditional Chinese, bring knowledge of the novels the world is built from, and act through a keyboard alone. It must keep a world model of its own, plan without a quest log or score, and fight turn-based battles against the AI of the game, over thousands of dependent steps with sparse reward. Every model plays under the same brief and control surface and is scored against events of the game itself. At the 60-minute budget the field of 11 models stays inside the opening: every model reaches the world map, 4 speak with the hermit, and one, `gpt-6-astra`, takes the compass, recruits a companion and fights, while none gains experience. The one milestone the random baseline reaches marks the floor a model has to clear. One frontier model plays the game beyond its opening, and even it is at its first fight with fourteen books to find. The ending lies tens of hours of coherent, self-directed play ahead. Four evaluations follow this one. Sessions of four to twenty-four hours bring the campaign milestones and a completion time within reach and make token cost per milestone comparable across models. The same models with web access measure how much a walkthrough adds over the brief alone. A human reference run under the same brief and budget, documented as Wei et al. (2025) recommend, calibrates the milestones against a first-time player. Once a model finishes, completion time turns the leaderboard into a speedrun.

AI USE STATEMENT

We used generative AI tools, coding agents built on frontier language models, throughout this work. They implemented the benchmark service, the replay scanner and the scripts that generate every

number, figure and table in this paper from the session records. They recovered and reformatted session records, supported the qualitative analysis of the replays, and proposed and refined parts of the measurement design. They searched for and formatted related work, and they drafted and edited the text of this paper under our direction. We have not used generative AI tools to generate synthetic data, to formulate or prove mathematical claims, or to propose hypotheses. The models evaluated in the benchmark are the object of study. We have reviewed all AI-assisted work. Every number in the paper is regenerated from the session records at each build, every typed claim about the records is verified by a consistency script, and the frames and readings behind the qualitative analysis were checked against the replays. We take responsibility for the final content of this work, including text, claims and artifacts produced with the aid of generative AI.

ETHICS STATEMENT

The benchmark runs a commercial 1996 game that is not redistributed: every host supplies its own copy, and the frames in this paper identify the interface under study, as Appendix A explains. The work involves no human participants, and the public catalogue carries the metrics, videos and keypress timelines of model sessions and no personal data.

REPRODUCIBILITY STATEMENT

The benchmark code, the brief the models receive, the replay scanner and the scripts that generate every number, figure and table in this paper will be released with the paper. The release includes the records, videos and keypress timelines of every session in the field. Appendix B lists the emulation and hosting parameters and the dates of the field, and Appendix C gives the reasons for the protocol choices. The game itself is not redistributable, as Appendix A states.

REFERENCES

- Anthropic. Claude’s extended thinking. Research post, section “Claude plays Pokémon”, <https://www.anthropic.com/research/visible-extended-thinking>, 2025.
- ARC Prize Foundation. ARC-AGI-3: A new challenge for frontier agentic intelligence. *arXiv preprint arXiv:2603.24621*, 2026.
- Axel Backlund and Lukas Petersson. Vending-Bench: A benchmark for long-term coherence of autonomous agents, 2025. URL <https://arxiv.org/abs/2502.15840>.
- Bahamut forum contributors. 原版金庸群俠傳攻略 (a complete walkthrough of the original game). <https://forum.gamer.com.tw/C.php?bsn=10859&snA=2778>, 2023.
- Marc G Bellemare, Yavar Naddaf, Joel Veness, and Michael Bowling. The Arcade Learning Environment: An evaluation platform for general agents. *Journal of Artificial Intelligence Research*, 47:253–279, 2013.
- Mingguang Chen, Licheng Wang, and Bo Qu. The horizon gap: Planning, memory, execution, training, and evaluation for long-horizon LLM agents, 2026. URL <https://arxiv.org/abs/2608.06663>.
- Gheorghe Comanici, Eric Bieber, Mike Schaeckermann, Ice Pasupat, Noveen Sachdeva, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities, 2025. URL <https://arxiv.org/abs/2507.06261>.
- Hafsteinn Einarsson. MazeEval: A benchmark for testing sequential decision-making in language models, 2025. URL <https://arxiv.org/abs/2507.20395>.
- Wei Fan, Xinjie Shen, Xudong Guo, Jianhong Tu, Yang Su, Yinger Zhang, Lianghao Deng, Fengyu Wang, Baohua Dong, Yangqiu Song, and Dayiheng Liu. E-Commerce Bench: Evaluating LLM agents on long-horizon autonomous business operation, 2026. URL <https://arxiv.org/abs/2608.30730>.
- Danijar Hafner. Benchmarking the spectrum of agent capabilities. In *International Conference on Learning Representations*, 2022.

- Lanxiang Hu, Mingjia Huo, Yuxuan Zhang, Haoyang Yu, Eric P Xing, Ion Stoica, Tajana Rosing, Haojian Jin, and Hao Zhang. *Imgame-Bench: How good are LLMs at playing games?* In *International Conference on Learning Representations*, 2026.
- Yuhan Jiang, Peng Luo, and Liqiu Meng. *Lost in Aggregation: A multi-scale diagnostic benchmark for LLM spatial navigation*, 2026. URL <https://arxiv.org/abs/2606.22219>.
- Tianjie Ju, Yueqing Sun, Zheng Wu, Wei Zhang, Yaqi Huo, Xi Su, Qi Gu, Xunliang Cai, Gongshen Liu, and Zhuosheng Zhang. *MineExplorer: Evaluating open-world exploration of MLLM agents in Minecraft*, 2026. URL <https://arxiv.org/abs/2605.30931>.
- Sayash Kapoor, Peter Kirgis, Andrew Schwartz, Stephan Rabanser, J. J. Allaire, Rishi Bommasani, Harry Coppock, Magda Dubois, Gillian K Hadfield, Andrew B. Hall, Sara Hooker, Seth Lazar, Steve Newman, Dimitris Papailiopoulos, Shoshannah Tekofsky, Helen Toner, Cozmin Ududec, and Arvind Narayanan. *Open-world evaluations for measuring frontier AI capabilities*, 2026. URL <https://arxiv.org/abs/2605.20520>.
- Seth Karten, Jake Grigsby, Tersoo Upaa Jr, Junik Bae, Seonghun Hong, Hyunyoung Jeong, Jaeyoon Jung, Kun Kerdtaisong, Gyungbo Kim, Hyeokgi Kim, Yujin Kim, Eunju Kwon, Dongyu Liu, Patrick Mariglia, Sangyeon Park, Benedikt Schink, Xianwei Shi, Anthony Sistilli, Joseph Twin, Arian Urdu, Matin Urdu, Qiao Wang, Ling Wu, Wenli Zhang, Kunsheng Zhou, Stephanie Milani, Kiran Vodrahalli, Amy Zhang, Fei Fang, Yuke Zhu, and Chi Jin. *The PokeAgent Challenge: Competitive and long-context learning at scale*, 2026a. URL <https://arxiv.org/abs/2603.15563>. NeurIPS 2025 Competition Track.
- Seth Karten, Joel Zhang, Tersoo Upaa Jr, Ruihong Feng, Wenzhe Li, Chengshuai Shi, Chi Jin, and Kiran Vodrahalli. *Continual harness: Online adaptation for self-improving foundation agents*, 2026b. URL <https://arxiv.org/abs/2605.09998>.
- Heinrich Küttler, Nantas Nardelli, Alexander Miller, Roberta Raileanu, Marco Selvatici, Edward Grefenstette, and Tim Rocktäschel. *The NetHack learning environment*. *Advances in Neural Information Processing Systems*, 33:7671–7684, 2020.
- Thomas Kwa, Ben West, Joel Becker, Amy Deng, Katharyn Garcia, Max Hasin, Sami Jawhar, Megan Kinniment, Nate Rush, Sydney Von Arx, Ryan Bloom, Thomas Broadley, Haoxing Du, Brian Goodrich, Nikola Jurkovic, Luke Harold Miles, Seraphina Nix, Tao Lin, Chris Painter, Neev Parikh, David Rein, Lucas Jun Koba Sato, Hjalmar Wijk, Daniel M. Ziegler, Elizabeth Barnes, and Lawrence Chan. *Measuring AI ability to complete long software tasks*, 2025. URL <https://arxiv.org/abs/2503.14499>. NeurIPS 2025.
- Wenhao Li, Wenwu Li, Chuyun Shen, Junjie Sheng, Zixiao Huang, Di Wu, Yun Hua, Wei Yin, Xiangfeng Wang, Hongyuan Zha, et al. *TextAtari: 100k frames game playing with language agents*. *arXiv preprint arXiv:2506.04098*, 2025.
- Ling Lin, Yang Bai, Heng Su, Congcong Zhu, Yaoming Wang, Yang Zhou, Huazhu Fu, and Jingrun Chen. *OODBench: Out-of-distribution benchmark for large vision-language models*, 2026.
- Qing’an Liu, Juntong Feng, Yuhao Wang, Xinzhe Han, Yujie Cheng, Yue Zhu, Haiwen Diao, Yunzhi Zhuge, and Huchuan Lu. *VISTA-Bench: Do vision-language models really understand visualized text as well as pure text?*, 2026.
- Haotian Luo, Huaisong Zhang, Xuelin Zhang, Haoyu Wang, Zeyu Qin, Wenjie Lu, Guozheng Ma, Haiying He, Yingsha Xie, Qiyang Zhou, Zixuan Hu, Hongze Mi, Yibo Wang, Naiqiang Tan, Hong Chen, Yi R. Fung, Chun Yuan, and Li Shen. *UltraHorizon: Benchmarking agent capabilities in ultra long-horizon scenarios*, 2025. URL <https://arxiv.org/abs/2509.21766>.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Belle-mare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharmarajan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. *Human-level control through deep reinforcement learning*. *Nature*, 518(7540):529–533, 2015.

- Junbo Niu, Yuanhong Zheng, Ziyang Miao, Hejun Dong, Chunjiang Ge, Hao Liang, Ma Lu, Bohan Zeng, Qiahao Zheng, Conghui He, and Wentao Zhang. Native visual understanding: Resolving resolution dilemmas in vision-language models, 2025.
- Mingyu Ouyang, Siyuan Hu, Kevin Qinghong Lin, Hwee Tou Ng, and Mike Zheng Shou. Game-World: Towards standardized and verifiable evaluation of multimodal game agents, 2026. URL <https://arxiv.org/abs/2604.07429>.
- Davide Paglieri, Bartłomiej Cupiał, Samuel Coward, Ulyana Piterbarg, Maciej Wołczyk, Akbir Khan, Eduardo Pignatelli, Łukasz Kuciński, Lerrel Pinto, Rob Fergus, et al. BALROG: Benchmarking agentic LLM and VLM reasoning on games. In *International Conference on Learning Representations*, 2025.
- Dongmin Park, Minkyu Kim, Beongjun Choi, Junhyuck Kim, Keon Lee, Jonghyun Lee, Inkyu Park, Byeong-Uk Lee, Jaeyoung Hwang, Jaewoo Ahn, et al. Orak: A foundational benchmark for training and evaluating LLM agents on diverse video games. In *International Conference on Learning Representations*, 2026.
- Long Phan, Mantas Mazeika, Andy Zou, and Dan Hendrycks. TextQuests: How good are LLMs at text-based video games? *arXiv preprint arXiv:2507.23701*, 2025.
- Pooyan Rahmazadehgervi, Logan Bolton, Mohammad Reza Taesiri, and Anh Totti Nguyen. Vision language models are blind: Failing to translate detailed visual features into words, 2024. URL <https://arxiv.org/abs/2407.06581>.
- scarsty and contributors. kys-cpp: a C++ reimplement of 金庸群俠傳. GitHub repository, <https://github.com/scarsty/kys-cpp>, 2026.
- Jiaqi Shao, Hanck Chen, Wei Zhang, Maxm Pan, and Bing Luo. Do agent benchmarks measure capability? Protocol validity in the age of agentic AI, 2026. URL <https://arxiv.org/abs/2607.22368>.
- Zhi Rui Tam, Ya-Ting Pai, Yen-Wei Lee, and Yun-Nung Chen. VisTW: Benchmarking vision-language models for Traditional Chinese in Taiwan, 2025. URL <https://arxiv.org/abs/2503.10427>.
- Muhammad Usama, Syeda Aishah Asim, Syed Bilal Ali, Syed Talal Wasim, and Umair Bin Mansoor. Analysing the robustness of vision-language models to common corruptions, 2025.
- Jiayu Wang, Yifei Ming, Zhenmei Shi, Vibhav Vineet, Xin Wang, Yixuan Li, and Neel Joshi. Is a picture worth a thousand words? Delving into spatial reasoning for vision language models, 2024. URL <https://arxiv.org/abs/2406.14852>. NeurIPS 2024.
- Archana Warriar, Dat Nguyen, Michelangelo Naim, Moksh Jain, Yichao Liang, Karen Schroeder, Cambridge Yang, Joshua B Tenenbaum, Sebastian Vollmer, Kevin Ellis, et al. Benchmarking world-model learning with environment-level queries. *arXiv preprint arXiv:2510.19788*, 2025.
- Nicholas R. Waytowich, Devin White, MD Sunbeam, and Vinicius G. Goecks. Atari-GPT: Benchmarking multimodal large language models as low-level policies in Atari games, 2024. URL <https://arxiv.org/abs/2408.15950>.
- Kevin L. Wei, Patricia Paskov, Sunishchal Dev, Michael J. Byun, Anka Reuel, Xavier Roberts-Gaal, Rachel Calcott, Evie Coxon, and Chinmay Deshpande. Position: Human baselines in model evaluations need rigor and transparency (with recommendations & reporting checklist). In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 82265–82325, 2025. URL <https://arxiv.org/abs/2506.13776>.
- Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, et al. OSWorld: Benchmarking multimodal agents for open-ended tasks in real computer environments. *Advances in Neural Information Processing Systems*, 37:52040–52094, 2024.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains, 2024. URL <https://arxiv.org/abs/2406.12045>.

Alex L Zhang, Thomas L Griffiths, Karthik R Narasimhan, and Ofir Press. VideoGameBench: Can vision-language models complete popular video games? *arXiv preprint arXiv:2505.18134*, 2025a.

Haoran Zhang, Chenhao Zhu, Sicong Guo, Hanzhe Guo, Haiming Li, and Donglin Yu. StarBench: A turn-based RPG benchmark for agentic multimodal decision-making and information seeking, 2025b. URL <https://arxiv.org/abs/2510.18483>.

Ruizhi Zhang, Ye Huang, Yuangang Pan, Chuanfu Shen, Zhilin Liu, Ting Xie, Haijun Lei, and Lixin Duan. Mastering PokeGym: Graph-guided multimodal evolution at test time, 2026a. URL <https://arxiv.org/abs/2604.08340>.

Yunbei Zhang, Janet Wang, Yingqiang Ge, Weijie Xu, Jihun Hamm, and Chandan K. Reddy. Stop comparing LLM agents without disclosing the harness, 2026b. URL <https://arxiv.org/abs/2605.23950>.

Yuxuan Zhu, Tengjun Jin, Yada Pruksachatkun, Andy Zhang, Shu Liu, Sasha Cui, Sayash Kapoor, Shayne Longpre, Kevin Meng, Rebecca Weiss, et al. Establishing best practices in building rigorous agentic benchmarks. *Advances in Neural Information Processing Systems*, 38, 2025.

A RELEASE AND LICENSING

The benchmark code is released under MIT. The emulation core (`dosbox_pure_libretro`) is GPLv2 and is loaded at runtime through the libretro C API. The bundled macOS core is built from `schellingb/dosbox-pure` at `7f6e8fb`, and the hosted service ships a Linux build pinned by image digest. The game is the 1996 commercial release, copyright Soft-World International (智冠科技) and Heluo Studio. It is not redistributable or tracked in the repository, and every host supplies its own copy. The public catalogue therefore carries only the metrics, videos and replay timelines we measured, and nothing that substitutes for the game data. The frames of the game in this paper identify the interface under study, in the same spirit as prior work showing frames of copyrighted games (Paglieri et al., 2025); if the rights holders object, they will be removed.

B INTERFACE AND SESSION DETAILS

Figure 5 shows the keyboard the API exposes. A host of 8 vCPUs and 16 GiB runs up to 24 concurrent sessions. Emulation is DOSBox Pure through libretro at a fixed 26800 cycles, the speed of a 486DX2-66 and period-correct for the game, with no sound card, so the observation is visual only. The published video runs at 8 times speed, or faster for long runs, with a strip naming the model, the action number, the keys held and the clock, beside a timeline file with every keypress and its hold. The sessions ran from 10 to 12 September 2026 under the model names the harness declared, the identifiers the providers expose, with the thinking level omitted since every session ran at the high level. Table 3 lists the effort of every model; the service sees no token counts, which are visible to the harness and the provider alone. The catalogue also holds 2 sessions of `gemini-3.6-flash`, a model with two later versions in the field, which the paper leaves out.

C DESIGN DECISIONS

An action accepts one key or a list of up to 100, because a menu path or a repeated step is one decision. A list costs one round trip and one settle wait for the whole list. The list does not loosen the coupling between perception and action. An action returns no picture, so a model that sends single keys without asking for the screen is as open-loop as one that sends twenty, and the coupling is set by how often a model looks. The cap bounds the worst case, every list is in the public timeline, and the most list-heavy model sent 22 keys per action on average. A single-key protocol would also change the action space the recorded field played under.




Figure 5: The keyboard the API exposes, with the nine keys the game responds to outside dialogue highlighted; any key advances a dialogue.

Table 3: Effort per model in the order of Figure 3; actions and seconds are medians over sessions.

Model	Sessions	Actions per session	Keys per action	Seconds between actions
gpt-6-astra	4	254	3.7	11
claude-opus5	2	303	6.5	12
gemini-3.8-flash	4	174	3.1	18
gpt-5.6-luna	5	155	6.9	11
gemini-3.7-flash	4	170	4.9	16
gpt-5.6-terra	6	126	11.8	12
claude-sonnet5	2	369	2.8	9
gpt-5.6-sol	4	332	3.1	8
claude-fable5.1	2	874	1.4	10
grok-4.6	2	472	1.4	5
glm-5.3-flash	5	215	1.8	14
random-baseline	2	2118	1.0	1

The observation is the frame the game draws, at its own resolution, since reading sixteen-pixel glyphs and an isometric scene is part of the task and the hermit milestone separates the models on that reading. A text channel or an upscaled frame would score a different task, and a fixed OCR component would put its own errors between the game and every model. Upscaling adds no information, and a model can resample its frame with the tools of its harness. The brief lists the on-screen terms in the Traditional Chinese the game draws.

The scores describe the 1996 release run as its own binary; a reimplementation such as kys-cpp (scarsty & contributors, 2026) has its own timing, text rendering and behaviour, so a track built on it would be a different benchmark. The memory reader anchors on one character name and confirms the two neighbouring records at the expected stride before it reads anything, so a copy with a different layout yields no reading, never a wrong one.

The baseline is a seeded random policy because it measures what the milestone structure gives away: the small first room puts the chest within reach of a random walk, and nothing after it is. A scripted policy would measure its author and a trained perception policy its training, while the leaderboard ranks frontier models under one harness. A human first-time run under the same brief and budget is part of the next evaluation. The 60-minute budget is sized to the opening, and the campaign budgets follow with the same milestones. With 2 to 6 sessions per model, the paper reports shares and spreads and does not claim that one model outranks another by a margin.

The readings are checked against one another wherever two records exist. Of the 10 crossings that carry both a save and a screen fingerprint, the two agree on 9, the save credits 1 crossing the fingerprint missed, and the fingerprint never appears without a save behind it. The crossing count read from the replay agrees with the count the service recorded on all 18 sessions that carry both. Over the 21 sessions with both a bag reading and a replay, every one of the 3 sessions whose bag holds the compass shows the coordinate line on the replay, and no other session does. Over the 30 sessions with a preserved save and a replay, every one of the 3 sessions whose save shows the scenes the hermit opens also shows his portrait on the replay. No frame without a panel scored above 0.79, and every session with an event has a frame above 0.90, so any threshold between the two gives the same readings.